

P Rakshith Reddy

Bangalore, India | rakshith488654@gmail.com | +91 8762210883

[LinkedIn/p-rakshith-reddy](#) | [GitHub/PRakshithReddy17](#) | [portfolio-three-rose-45.vercel.app](#)

Summary

Engineering student in Computer Science (AI & ML) at Jain University, specializing in building and deploying machine learning systems and full-stack web applications. Experienced with Python, PyTorch, FastAPI, and React across 4+ production projects. Seeking ML internship or full-time software engineering roles.

Education

Jain University, Bangalore, India

Expected: Aug 2028

Bachelor of Technology, Computer Science (AI & Machine Learning)

CGPA: 8.79 / 10.0

Relevant Coursework: Machine Learning, NLP, Computer Vision, Data Structures, Algorithms, Web Development

Technical Skills

Languages: Python, TypeScript, JavaScript, C++

AI/ML: PyTorch, TensorFlow, Scikit-learn, Hugging Face, FAISS, Transformers, RAG, LLM Integration, Prompt Engineering, Pandas, NumPy

Frameworks: React, Next.js, FastAPI, Streamlit, Tailwind CSS

Tools & Databases: Docker, Git, Vercel, LangChain, OpenRouter API, Scikit-build

Projects

AI Resume Analyzer + Job Matcher [[Link](#) | [Demo](#)]

2024–2025

Streamlit Web Application

- Developed NLP-based resume-job matching system using TF-IDF vectorization and cosine similarity
- Designed hybrid scoring algorithm (60% skill match, 40% similarity) achieving up to 91% match accuracy
- Built interactive Streamlit dashboard with real-time match scoring and skill-gap insights (1–2 sec latency)

Multilingual Toxic Filter [[Link](#)]

2025

Full-stack NLP Application

- Built multilingual toxic comment detection system using pretrained BERT-based (Toxic-BERT) transformer model, enabling classification across 100+ languages
- Engineered FastAPI-based inference pipeline achieving 100–300 ms latency per request for real-time moderation
- Implemented adversarial input handling including obfuscation detection and Unicode script validation to prevent bypass attacks

AI Career Counselor Chatbot [[Link](#) | [Demo](#)]

2025

Full-stack AI Application

- Developed full-stack AI chatbot using Next.js and FastAPI for personalized career guidance
- Integrated GPT-3.5 Turbo via OpenRouter API with prompt-engineered system role for structured responses
- Designed stateless architecture with dynamic prompt construction for context-aware replies
- Achieved response latency of 1–3 seconds; deployed on Vercel/Render with real-time interaction

AI Research Summarizer [[Link](#) | [Demo](#)]

2025–2026

RAG Pipeline with Vector Indexing

- Built RAG-based research summarization system using FAISS and Sentence-Transformers embeddings
- Implemented chunking (1000 chars, 200 overlap) and top-k retrieval (k=5) for relevant context extraction
- Optimized performance using lightweight MiniLM embeddings and FAISS similarity search
- Integrated local LLM (phi3:mini) achieving 2–5 sec response time for PDF-based queries

Certifications

[ML with Python \(IBM\)](#) — [Machine Learning in Python](#)

[Data Engineering \(ETL\)](#) — [Operating Systems](#)

[Quantum Mechanics](#)

Languages & Interests

Languages: English (Fluent), Hindi, Telugu, Kannada

Interests: Open Source Contribution, Technical Content Writing, Competitive Programming, Cricket